

CS344 LAB 1 Assignment

Yamini Sunil Tasare 220123071

Vishal Choudhary 220123070

Sneh Dadhania 220123062

Subhrajyoti Kunda Roy 220123064

Drive link for code of Part 0B - 1 - [code](#)

PART 0A-1

Exercise 1

```
C ex1.c > main(int, char **)
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      int x = 1;
6      printf("Hello x = %d\n", x);
7
8      // Inline assembly |
9      __asm__("incl %0" : "=r"(x) : "0"(x));
10
11     printf("Hello x = %d after increment\n", x);
12
13     if(x == 2){
14         printf("OK\n");
15     }
16     else{
17         printf("ERROR\n");
18     }
19 }
20
```

PROBLEMS OUTPUT TERMINAL PORTS JUPYTER

> **TERMINAL**

```
PS E:\why> gcc ex1.c
PS E:\why> ./a.exe
Hello x = 1
Hello x = 2 after increment
OK
PS E:\why> 
```

Exercise 2

```
(gdb) target remote localhost:26000
Remote debugging using localhost:26000
mycpu () at proc.c:48
48      for (i = 0; i < ncpu; ++i) {
(gdb) si
0x80103907      48      for (i = 0; i < ncpu; ++i) {
(gdb) si
0x80103909      48      for (i = 0; i < ncpu; ++i) {
(gdb) si
0x8010390b      48      for (i = 0; i < ncpu; ++i) {
(gdb) si
0x8010390d      48      for (i = 0; i < ncpu; ++i) {
(gdb) si
49      if (cpus[i].apicid == apicid)
(gdb) si
0x8010391d      49      if (cpus[i].apicid == apicid)
(gdb) si
0x80103924      49      if (cpus[i].apicid == apicid)
(gdb) si
0x80103926      49      if (cpus[i].apicid == apicid)
(gdb) si
50      return &cpus[i];
(gdb) si
0x8010392b      50      return &cpus[i];
(gdb)
```

Commands

```
$ cd xv6-public
```

```
$ gdb kernel
```

// the above command launches GDB and loads the debugging symbols for the kernel executable. This allows us to debug the xv6 kernel.

```
(gdb) target remote localhost:26000
```

What the code seems to be doing :

Iterating through an array of CPU structures(cpus)

Checking if the `apicid` of each CPU matches a given `apicid`, if so, returning a pointer to that CPU structure.

This function appears to be part of the kernel's CPU management, likely used to find the data structure representing the current CPU based on its `APIC ID` (Advanced Programmable Interrupt Controller ID).

While `si` isn't directly accessing line 48 of `proc.c`, it's allowing us to step through the machine instructions that correspond to that line and subsequent lines of source code, instruction by instruction. Each time we enter `si`, GDB executes one assembly instruction.

After each `si`, GDB shows:

The memory address of the next instruction (e.g., `0x80103907`, `0x80103909`, etc.)

The corresponding line number in the source code.

The source code line itself.

Exercise 3

Setting breakpoint at 0x7c00 :

```
(gdb) source .gdbinit
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is set to "i8086".
[f000:ffff] 0xffff0: jmp $0x3630,$0xf000e05b
0x0000ffff in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration
of GDB. Attempting to continue with the default i8086 settings.

(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) si
[f000:e062] 0xfe062: jne 0xd241d0b0
0x0000e062 in ?? ()
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) si
[ 0:7c01] => 0x7c01: xor %eax,%eax
0x00007c01 in ?? ()
(gdb) si
[ 0:7c03] => 0x7c03: mov %eax,%ds
0x00007c03 in ?? ()
(gdb) |
```

a)

The processor starts executing 32-bit code -

Corresponding code in [bootblock.asm](#) ->

```
//PAGEBREAK!
# Complete the transition to 32-bit protected mode by using a long
jmp
# to reload %cs and %eip. The segment descriptors are set up with
no
# translation, so that the mapping is still the identity mapping.
ljmp $(SEG_KCODE<<3), $start32
7c2c: ea .byte 0xea
7c2d: 31 7c 08 00 xor %edi,0x0(%eax,%ecx,1)
```

Corresponding code in [bootasm.S](#) ->

```
0x7c29: mov %eax,%cr0
0x7c2c: ljmp $0xb866,$0x87c31
0x7c33: adc %al,(%eax)
0x7c35: mov %eax,%ds
0x7c37: mov %eax,%eax
```

b)

The last instruction of the boot loader executed

In [bootblock.asm](#) ->

```
7d7f: 72 15 jb 7d96 <bootmain+0x59>
entry();
7d81: ff 15 18 00 01 00 call *0x10018
}
7d87: 8d 65 f4 lea -0xc(%ebp),%esp
```

Corresponding code in `bootmain.c` ->

```
// Call the entry point from the ELF header.
// Does not return!
entry = (void(*)(void))(elf->entry);
entry();
}
```

The address 0x10018 is called. (First instruction of Kernel)

First instruction in kernel -> `movl %cr,%eax`

```
// bootmain() loads an ELF kernel image from the disk starting at
// sector 1 and then jumps to the kernel entry routine.
```

c)

```
35  ph = (struct proghdr*)((uchar*)elf + elf->phoff);
36  eph = ph + elf->phnum;
37  for(; ph < eph; ph++){
38      pa = (uchar*)ph->paddr;
39      readseg(pa, ph->filesz, ph->off);
40      if(ph->memsz > ph->filesz)
41          stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
42  }
```

`ph` is a pointer to the program header, `elf` contains the number of sectors required to fetch the entire kernel, `eph` is the pointer to the last sector.

The for loop loads each sector from `ph` upto `eph` incrementing `ph` at each iteration. This information is present in `elf` headers.

Exercise 5

In `bootblock.asm`, the `lgdtl (%esi)` will break or do the wrong thing if we set the boot loader's link to the wrong address.

```
# Switch from real to protected mode. Use a bootstrap GDT that makes
# virtual addresses map directly to physical addresses so that the
# effective memory map doesn't change during the transition.
lgdt    gdt_desc
7c1d:    0f 01 16                lgdtl    (%esi)
7c20:    78 7c                js      7c9e <readsect+0x12>
```

Changed the bootloader's link address to `0x8c00` from `0x7c00`.

```

103 bootblock: bootasm.S bootmain.c
104 $(CC) $(CFLAGS) -fno-pic -O -nostdinc -I. -c bootmain.c
105 $(CC) $(CFLAGS) -fno-pic -nostdinc -I. -c bootasm.S
106 $(LD) $(LDFLAGS) -N -e start -Ttext 0x8C00 -o bootblock.o bootasm.o bootmain.o
107 $(OBJDUMP) -S bootblock.o > bootblock.asm
108 $(OBJCOPY) -S -O binary -j .text bootblock.o bootblock
109 ./sign.pl bootblock

```

Further, ran the 'make clean' and 'make' commands

```

yamini@yamini:~/xv6-public$ make clean
rm -f *.tex *.dvi *.idx *.aux *.log *.ind *.ilg \
*.o *.d *.asm *.sym vectors.S bootblock entryother \
initcode initcode.out kernel xv6.img fs.img kernelmemfs \
xv6memfs.img mkfs .gdbinit \
_cat _echo _forktest _grep _init _kill _ln _ls _mkdir _rm _sh _stressfs _use
rtests _wc _zombie
yamini@yamini:~/xv6-public$ make
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -f
no-pic -O -nostdinc -I. -c bootmain.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -f
no-pic -nostdinc -I. -c bootasm.S
ld -m elf_i386 -N -e start -Ttext 0x8C00 -o bootblock.o bootasm.o bootmai
n.o
objdump -S bootblock.o > bootblock.asm
objcopy -S -O binary -j .text bootblock.o bootblock
./sign.pl bootblock
boot block is 451 bytes (max 510)
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o bio.o bio.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o console.o console.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o exec.o exec.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o file.o file.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o fs.o fs.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o ide.o ide.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o ioapic.o ioapic.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie

```

After running these commands, we observe that the address in the bootloader changes but the ip is still 7c00.

```

yamini@yamini:~/xv6-public$ cat bootblock.asm
bootblock.o:      file format elf32-i386

Disassembly of section .text:

00008c00 <start>:
# with %cs=0 %ip=7c00.

.code16                                # Assemble for 16-bit mode
.globl start
start:
cli                                     # BIOS enabled interrupts; disable
    8c00:      fa                    cli

# Zero data segment registers DS, ES, and SS.
xorw  %ax,%ax                          # Set %ax to zero
    8c01:      31 c0                xor  %eax,%eax
movw  %ax,%ds                          # -> Data Segment
    8c03:      8e d8                mov  %eax,%ds
movw  %ax,%es                          # -> Extra Segment
    8c05:      8e c0                mov  %eax,%es
movw  %ax,%ss                          # -> Stack Segment
    8c07:      8e d0                mov  %eax,%ss

00008c09 <seta20.1>:

# Physical address line A20 is tied to zero so that the first PCs
# with 2 MB would run software that assumed 1 MB. Undo that.
seta20.1:
inb   $0x64,%al                        # Wait for not busy
    8c09:      e4 64                in    $0x64,%al
testb $0x2,%al
    8c0b:      a8 02                test  $0x2,%al
jnz   seta20.1
    8c0d:      75 fa                jne   8c09 <seta20.1>

```

Exercise 6

```
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) x/8i 0x00100000
0x100000: add %al, (%eax)
0x100002: add %al, (%eax)
0x100004: add %al, (%eax)
0x100006: add %al, (%eax)
0x100008: add %al, (%eax)
0x10000a: add %al, (%eax)
0x10000c: add %al, (%eax)
0x10000e: add %al, (%eax)
(gdb) b *0x7d81
Breakpoint 2 at 0x7d81
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x7d81: call *0x10018

Thread 1 hit Breakpoint 2, 0x00007d81 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x83e0200f
0x100010: 0x220f10c8 0x9000b8e0 0x220f0010 0xc0200fd8
```

8 words of instruction at `0x00100000` at the point when BIOS enters the boot loader and 8 words of instruction at `0x00100000` at the point when the boot loader enters the kernel are different as when the BIOS enters and loads the boot loader, then it just loads it in memory location between `0x7C00` and `0x7DFF` due to which all the 8 words of instructions are zero at `0x00100000`. But before the boot loader enters the kernel, it already has performed the 16 to 32 bit transition and setting up of stack which leads to new instructions at address `0x00100000`

The address `0x7DFF` comes from the memory range allocated for the boot sector. When the BIOS loads the boot sector into memory, it places it at the address `0x7C00`, which is the starting address.

The boot sector typically occupies 512 bytes of memory, so it spans from `0x7C00` to `0x7DFF`. The `0x7DFF` is simply the last byte of this 512-byte space, calculated as:

Start address: `0x7C00`, Size: 512 bytes, End address: $0x7C00 + 0x0200 - 1 = 0x7DFF$.

```
yamini@yamini:~/xv6-public$ cat bootmain.c
// Boot loader.
//
// Part of the boot block, along with bootasm.S, which calls bootmain().
// bootasm.S has put the processor into protected 32-bit mode.
// bootmain() loads an ELF kernel image from the disk starting at
// sector 1 and then jumps to the kernel entry routine.

#include "types.h"
#include "elf.h"
#include "x86.h"
#include "memlayout.h"

#define SECTSIZE 512
```

PART 0B-1

Exercise 1

To define a system call in xv6, we made the following changes

1) syscall.h

We added a new system call `#define SYS_draw` at 22nd position as 21 positions were already occupied by the inbuilt system calls in syscall.h

2) syscall.c

We added a pointer `[SYS_draw] sys_draw` to system call at 22nd position in syscall.c file in order to add our custom system call.

Then a function prototype `extern int sys_draw(void);` is added in syscall.h file which will be called by system call number 22.

3) sysproc.c

System call function is implemented in sysproc.c

```
95 int sys_draw(void){
96     char *a =
97     "   .-. \n"
98     "  |o_o | \n"
99     "  |:_/  | \n"
100    " //   \ \  \ \ \n"
101    " (|    | ) \n"
102    " /'\_/_/^\_ \ \ \n"
103    " \\/_/\\_/_/ \n";
104
105    void* buf;
106    int size;
107
108    if(argint(1,&size) < 0){
109        return -1;
110    }
111
112    if(argptr(0,(char*)&buf,size) < 0){
113        return -1;
114    }
115    int p = 0;
116    for(p = 0; a[p] != '\0'; p++);
117
118    if(p > size) p = -1;
119    else {
120        for(int i = 0; i < p; i++) ((char*)buf)[i] = a[i];
121    }
122    // returning size of wolfie
123    // -1 is returned if buffer created is not enough to hold wolfie otherwise wolfie size is returned
124    return p;
125 }
```

4) usys.S

For creating an interface for your user program to access system call we added the following line in usys.S

`SYSCALL(draw)`

5) user.h

We added the following function that the user program will be calling in user.h,

```
int draw(void *buf, uint size);
```

Call to the above function from the user program will be simply mapped to system call number 22 which is defined as `SYS_draw` preprocessor directive. The system knows what exactly this system calls and how to handle it.

Exercise 2

We save a C program named `testfile.c` inside the source code directory of xv6 operating system.

Then we edit the MakeFile and added below changes in MakeFile, Under the value `UPROGS` (present at line 168), we added `_testfile\` at the end of `UPROGS` value and Under the value `EXTRA` (present at line 251), we added `testfile.c\`.

Then we run:

'Make clean'

'Make'

'Make qemu'

'ls'

'testfile'.

```
$ ls
.          1 1 512
..         1 1 512
README    2 2 2286
cat       2 3 15472
echo      2 4 14348
forktest  2 5 8800
grep      2 6 18316
init      2 7 14972
kill      2 8 14436
ln        2 9 14332
ls        2 10 16904
mkdir     2 11 14456
rm        2 12 14436
sh        2 13 28496
stressfs  2 14 15368
usertests 2 15 62868
wc        2 16 15892
zombie    2 17 14016
testfile  2 18 14396
console   3 19 0
$ testfile
Penguin Size = 75 B
  .--.
  |o_o |
  |:/_|
 //  \ \
(|  |  |)
/\'  _/\'
 \--/\--/

pid 4 testfile: trap 14 err 5 on cpu 0 eip 0xffffffff addr 0xffffffff--kill proc
$ |
```